

Excellent. Now we move into **real SQL engineering territory**.

This is a **deep dive into query optimization**, not surface-level advice.

We'll cover:

1. How MySQL actually executes queries
2. How to read `EXPLAIN` properly
3. Index internals (B+Tree)
4. Composite indexes & leftmost rule
5. Covering indexes
6. Join optimization
7. Subquery vs join performance
8. Window function cost
9. Pagination optimization
10. Real optimization workflow

We'll use the `advanced_db` dataset created earlier.



1

How MySQL Executes a Query

When you run:

```
SELECT * FROM employees WHERE salary > 80000;
```

Internally MySQL does:

1. Parse SQL
2. Check syntax
3. Optimize (choose execution plan)
4. Execute

The optimizer chooses:

- Which index to use
 - Join order
 - Access method (index scan vs full scan)
-



2 FULL TABLE SCAN (Worst Case)

Without index:

```
SELECT * FROM employees WHERE salary > 80000;
```

MySQL checks **every row**.

This is O(n).

For 10 rows → fine

For 10 million rows → disaster



3 USING EXPLAIN (Most Important Tool)

```
EXPLAIN SELECT * FROM employees WHERE salary > 80000;
```

Key columns to understand:

| Column | Meaning |
|--------------|------------------------|
| type | Access method |
| possible_key | Available indexes |
| s | |
| key | Index actually used |
| rows | Estimated rows scanned |
| Extra | Extra info |

Access Types (Ordered Best → Worst)

| type | Meaning |
|--------|--------------|
| const | Single row |
| eq_ref | Perfect join |

ref Indexed lookup

range Range scan

index Full index scan

ALL Full table scan

If you see:

```
type = ALL
```

You have a performance problem.



4

ADD INDEX & TEST

```
CREATE INDEX idx_salary ON employees(salary);
```

Now run:

```
EXPLAIN SELECT * FROM employees WHERE salary > 80000;
```

Expected:

```
type = range  
key = idx_salary
```

Now MySQL uses **B+Tree index range scan**.



5

How B+Tree Index Works

Index stores:

```
salary → primary key pointer
```

Sorted order:

60000
70000
75000
80000
90000

Binary search is used → O(log n)

Then rows are fetched.



6 Composite Index (CRITICAL)

Suppose query:

```
SELECT * FROM employees
WHERE dept_id = 1 AND salary > 70000;
```

Better index:

```
CREATE INDEX idx_dept_salary
ON employees(dept_id, salary);
```

This is powerful because:

MySQL uses **leftmost prefix rule**.

Leftmost Rule

Index (`dept_id, salary`) works for:

- ✓ WHERE `dept_id = ?`
- ✓ WHERE `dept_id = ? AND salary > ?`
- ✗ WHERE `salary > ? alone`

Because `dept_id` is first column.

7 Covering Index (Advanced Optimization)

If query is:

```
SELECT dept_id, salary  
FROM employees  
WHERE dept_id = 1;
```

If index is:

```
CREATE INDEX idx_cover ON employees(dept_id, salary);
```

MySQL may not even touch the table.

Because:

All needed columns are inside index.

EXPLAIN shows:

Using index

That is extremely fast.

8 JOIN OPTIMIZATION

Bad:

```
SELECT *  
FROM orders o  
JOIN customers c ON o.customer_id = c.customer_id;
```

If no index on:

`orders.customer_id`

MySQL may scan entire orders table.

Fix:

```
CREATE INDEX idx_orders_customer  
ON orders(customer_id);
```

Foreign key columns should almost always be indexed.



9

JOIN ORDER MATTERS

MySQL optimizer decides join order.

Example:

```
SELECT *  
FROM employees e  
JOIN departments d ON e.dept_id = d.dept_id  
WHERE d.dept_name = 'Engineering';
```

If index on:

```
departments.dept_name
```

MySQL first filters departments → then joins employees.



10

Subquery vs Join

Slow:

```
SELECT *  
FROM employees  
WHERE dept_id IN  
(SELECT dept_id FROM departments WHERE dept_name = 'Engineering');
```

Better:

```
SELECT e.*  
FROM employees e
```

```
JOIN departments d  
ON e.dept_id = d.dept_id  
WHERE d.dept_name = 'Engineering';
```

Joins usually optimize better.



11 Avoid SELECT *

Bad:

```
SELECT * FROM employees;
```

Why?

- Transfers unnecessary columns
- Prevents covering index usage

Better:

```
SELECT emp_id, first_name, salary  
FROM employees;
```



12 Pagination Optimization

Bad:

```
SELECT * FROM employees  
LIMIT 100000, 10;
```

This skips 100k rows → slow.

Better approach:

Use indexed column:

```
SELECT *
```

```
FROM employees  
WHERE emp_id > 100000  
LIMIT 10;
```

This is keyset pagination.



13 ORDER BY Optimization

Bad:

```
SELECT * FROM employees  
ORDER BY salary;
```

If no index on salary → filesort.

Fix:

```
CREATE INDEX idx_salary ON employees(salary);
```

Then EXPLAIN should show:

Using index



14 Avoid Functions on Indexed Columns

Bad:

```
SELECT * FROM employees  
WHERE YEAR(hire_date) = 2023;
```

This disables index.

Better:

```
WHERE hire_date BETWEEN '2023-01-01' AND '2023-12-31';
```



15 Window Functions Cost

```
SELECT first_name,  
       RANK() OVER (ORDER BY salary DESC)  
FROM employees;
```

Window functions require sorting.

Large datasets → heavy memory usage.

Index on salary helps but still sorting happens.



16 Transaction & Lock Optimization

Long transactions → lock contention.

Bad:

```
START TRANSACTION;  
SELECT ...  
(wait 5 minutes)  
UPDATE ...  
COMMIT;
```

Keep transactions short.



17 Real Optimization Workflow

When query is slow:

Run:

```
EXPLAIN query;
```

1.

2. Check:
 - type = ALL? 
 - rows large?
 - no key used?
 3. Add appropriate index
 4. Avoid:
 - functions on indexed columns
 - OR across multiple columns (may disable index)
 5. Re-test EXPLAIN
-



18 Index Strategy Best Practices

| Situation | Index Strategy |
|----------------------------------|------------------------|
| Foreign keys | Always index |
| WHERE column | Index |
| WHERE + ORDER BY | Composite index |
| Low cardinality (gender, status) | Avoid indexing |
| Many updates | Avoid too many indexes |



19 Cardinality Matters

If column has only:

- 2 values (status: active/inactive)

Index not helpful.

If column has:

- Millions of distinct values (email, emp_id)

Very useful.



20 Performance Interview Questions

- What is covering index?
 - What is leftmost prefix rule?
 - Why avoid function on indexed column?
 - What is filesort?
 - What is range scan?
 - What is keyset pagination?
 - What is index selectivity?
-



Practice Optimization Exercises

Run these and optimize:

1.

```
SELECT * FROM orders WHERE status='Completed';
```

2.

```
SELECT * FROM employees WHERE salary > 70000 ORDER BY hire_date;
```

3.

```
SELECT c.customer_name, SUM(o.total_amount)
FROM customers c
JOIN orders o ON c.customer_id=o.customer_id
```

```
GROUP BY c.customer_id;
```

Add appropriate indexes and compare EXPLAIN.

Final Key Principle

Indexes speed up reads

Indexes slow down writes

Balance is required.